

OmniDash - Tech Spec (DRAFT)

Omni Dash

Technical Specification — YouTube MVP & Phased Platform Expansion

Prepared by: Second Life Software

Prepared for: Omni Dash

Version: 1.1

Document Type: Authoritative Technical Specification

Status: Implementation-Ready

1. Purpose of This Document

This document defines the **exact technical implementation** for the Omni Dash **YouTube-only MVP** and its phased expansion across additional platforms.

It serves as:

- The **single source of truth** for engineering execution
- A shared alignment artifact between Omni Dash and Second Life Software
- A technical appendix supporting the Statement of Work (SOW)

All decisions herein are **final, intentional, and scoped** to optimize for speed, safety, reliability, and cost efficiency.

2. MVP Product Scope (YouTube Only)

In Scope (Phase 1 MVP)

- Responsive **web application**
- AI-assisted and AI-autonomous comment engagement (YouTube only)

- **YouTube platform integration (required and sole platform for MVP)**
- Controlled rollout to **100–200 early users**
- Paid SaaS subscriptions (Stripe)
- Multi-layer AI guardrails with full auditability
- **Firebase Authentication** (creator login + onboarding)

Explicitly Out of Scope (Phase 1 MVP)

- Native mobile applications
 - Direct messages (DMs)
 - Enterprise workflows
 - Instagram / X / any second platform integration
 - More than one platform in MVP
 - Unrestricted or unsupervised AI autonomy
 - **LoRA / model training / fine-tuning**
-

3. Technology Stack (Final Decisions)

Frontend

- **React + TypeScript** (Next.js)
- Tailwind CSS

Authentication

- **Firebase Authentication**
 - Email/password (baseline)
 - Optional: Google OAuth as a follow-on if required by Omni Dash
- Firebase used for authentication only; user profile and application data remain in Postgres.

Backend API

- **Python FastAPI**
- Async-first architecture
- Stateless service design

Background Processing

- Python worker services
- Queue-based job execution
- Deterministic retry and backoff logic

Database

- **PostgreSQL**
- Primary relational datastore (Cloud SQL)

Caching & Rate Limiting

- **Redis**
- Used for:
 - reply deduplication
 - rate limiting
 - short-lived state locks

AI Inference

- Open-source LLM
 - llama 3.1
- Served via **vLLM**
- Dedicated GPU inference service

Hosting Provider

- **Google Cloud Platform (GCP)**

Deployment Model

- Web & API: **Cloud Run**
- Database: **Cloud SQL (Postgres)**
- Cache: Redis (managed or VM-hosted)
- GPU inference: **Compute Engine GPU VM**
- Object storage: **Cloud Storage**

This architecture isolates GPU cost while allowing all other services to scale to zero.

4. AI Model Selection

MVP Model (Phase 1)

- **Llama 3.1 8B Instruct**
- Quantized (4-bit or 8-bit)
- Temperature and token limits enforced

Model Serving

- **vLLM** inference server
 - Request batching enabled
 - Structured output enforced (JSON schema)
-

5. System Architecture Overview

Core Services

1. Web Application

- Creator dashboard
- Unified inbox
- AI approval workflows
- Analytics

- Subscription management

2. Authentication Service

- Firebase Authentication for creator identity
- Backend verifies Firebase JWT and maps to internal Omni Dash user records

3. API Service

- Authentication & authorization
- Platform orchestration
- Guardrail enforcement
- Data access layer

4. Worker Services

- Comment ingestion
- AI generation jobs
- Reply posting

5. AI Inference Service

- GPU-hosted LLM endpoint

6. Admin Console

- User management
- Kill switches
- Platform connection status

6. Platform Connector Architecture (YouTube Only in MVP)

Each platform is implemented as a **modular connector** with a standardized interface:

- OAuth authentication

- Comment ingestion
- Normalization to internal schema
- Reply posting
- Rate-limit handling

Phase 1 MVP includes the YouTube connector only.

Instagram and X connectors are explicitly deferred to later phases.

7. Phased Delivery Plan

Phase 0 — Core Platform Foundation (Included in Phase 1 MVP Execution)

Objective: Establish platform-agnostic infrastructure and AI orchestration (even though only YouTube is implemented in MVP).

Deliverables

- Firebase Authentication integration (creator login + session)
 - User/team model (creator + optional team members)
 - Stripe subscription integration (paid access gating)
 - Unified comment inbox foundation (platform-neutral internal schema)
 - AI orchestration pipeline (generate → validate → approve/post)
 - Guardrails framework (v1)
 - Audit logging system
 - Admin controls
 - Observability & logging
-

Phase 1 — YouTube MVP (Scope-Committed)

Objective: Deliver production-ready MVP functionality **for YouTube only**, including Assisted and Autonomous AI modes with full guardrails and auditability.

Phase 1 Assumptions

- AI development tools (Copilot / Cursor / ChatGPT) may be used to accelerate implementation, but all architecture and safety decisions remain human-owned.
 - **No LoRA / training / fine-tuning** in Phase 1.
 - **One platform only:** YouTube.
 - Web-first only.
 - Production-ready (reliable and auditable), not demo-only.
-

Phase 1 Work Breakdown (No Hours)

1) Product & Technical Planning

Deliverables

- Finalize architecture and service boundaries (Web/API/Workers/Inference)
 - Finalize data model (users, comments, threads, replies, audit logs, rulesets)
 - Finalize guardrail policy thresholds and enforcement ordering
 - Formalize platform connector contract (implemented for YouTube in MVP)
-

2) Frontend (Next.js Web App)

Deliverables

- Authentication + onboarding flows (Firebase Auth)
- Creator dashboard
- Unified inbox UI (threaded comments, filtering, search)
- Assisted Mode UX:
 - generate suggestions

- approve/edit
 - post
 - Autonomous Mode controls:
 - opt-in activation
 - reply limits
 - active hours
 - kill switch
 - Analytics views (MVP)
 - Production UI readiness:
 - error states
 - loading states
 - guardrail-block messaging
-

3) Backend API (FastAPI)

Deliverables

- Firebase JWT verification + internal user mapping
 - Core data models + API endpoints
 - Comment lifecycle state machine (ingest → eligible → draft → validate → approve/post → log)
 - Creator voice profile system:
 - tone, length, emoji rules, phrase bans, examples
 - Guardrail enforcement module:
 - eligibility, behavioral, safety checks
 - Assisted Mode orchestration endpoints
 - Autonomous Mode orchestration endpoints
 - Analytics aggregation endpoints (MVP-level)
-

4) YouTube Integration

Deliverables

- OAuth connection flow
 - Comment ingestion:
 - thread handling
 - pagination
 - polling cadence
 - Reply posting
 - Rate limit handling + retries
 - Integration testing + edge-case handling
-

5) AI Inference & Prompting

Deliverables

- GPU inference service setup (vLLM + Llama 3.1 8B Instruct)
 - Prompt templates:
 - system rules
 - voice profile injection
 - few-shot examples
 - structured JSON output
 - Voice example handling:
 - capture approved replies
 - reuse as examples
 - Safety validation:
 - post-generation checks
 - optional “judge” evaluation pass (MVP level)
 - Prompt testing across known risky inputs and edge cases
-

6) Guardrails & Safety (Cross-cutting)

Deliverables

- Reply deduplication (no multiple replies per comment)
 - Rate limiting (per user + global caps)
 - Active hours enforcement
 - PII detection and blocking
 - Toxicity / unsafe input filtering
 - Audit logging + traceability:
 - why a reply was allowed/blocked
 - what rules were applied
 - what mode posted it (assisted/autonomous)
-

7) Infrastructure & Deployment (GCP)

Deliverables

- Cloud Run deployment for web + API
 - Cloud SQL provisioning and migrations
 - Redis provisioning
 - GPU VM provisioning for inference service
 - Secrets management via GCP Secret Manager
 - IAM / networking configuration
 - Observability:
 - structured logs
 - error reporting
 - basic alerting
-

8) QA, Stabilization & Delivery

Deliverables

- Integration testing across:
 - YouTube ingest
 - reply posting
 - assisted approvals
 - autonomous posting flows
 - Edge-case handling and bug fixes
 - Deployment hardening and rollback procedures
 - Client handoff:
 - walkthrough
 - runbook
 - admin training (kill switch, audit logs, platform connection troubleshooting)
-

Phase 2 — Instagram Integration (Deferred)

Objective: Add second platform using the existing connector framework.

Scope

- OAuth connection
 - Comment ingestion
 - Reply posting
 - Shared inbox & analytics
-

Phase 3 — X (Twitter) Integration (Deferred)

Objective: Add third platform post-MVP.

Scope

- Mentions/replies ingestion

- Reply posting
 - Enhanced throttling due to platform sensitivity
-

8. Guardrail System (Mandatory)

Omni Dash implements a **five-layer guardrail system** enforced deterministically by application logic.

Guardrail Layers

1. Eligibility
 2. Behavioral
 3. Voice & Content
 4. Safety & Compliance
 5. Control & Auditability
-

Guardrail Matrix

Layer	Description	Enforcement
Eligibility	Deduplication, topic blocks, toxic input	Code + classifiers
Behavioral	Rate limits, active hours, cooldowns	Code
Voice	Tone, length, banned phrases	Prompt constraints
Safety	PII, hate, claims, policy	Filters + classifiers
Control	Opt-in, kill switch, audit logs	UX + code

Mandatory Phase 1 MVP Guardrails

- Reply deduplication (no multiple replies per comment)
- Max replies per hour/day
- Active hours enforcement
- Autonomous Mode opt-in
- Kill switch

- Input & output safety checks
 - Full audit logging
-

9. Comment Lifecycle State Machine

1. Comment ingested
 2. Eligibility evaluated
 3. AI reply generated
 4. Post-generation safety validation
 5. Assisted approval or autonomous post
 6. Audit logged
 7. Analytics updated
-

10. Creator Voice System

Creators configure a **Voice Profile** consisting of:

- Tone
- Length preferences
- Emoji usage rules
- Phrase bans
- Example replies

Voice profiles are applied at generation time via structured prompts and example injection.

11. LoRA / Model Tuning Strategy (Out of Scope for Phase 1)

Phase X MVP Position

- **No model training in MVP**

- Prompting + voice profiles only

Post-MVP (Phase 4+)

LoRA tuning will be introduced once:

- 300–1,000+ approved replies exist
- Measurable quality improvements are required
- Cost or latency optimization becomes necessary

LoRA adapters will be:

- Versioned
 - Optional
 - Reversible
-

12. Security & Reliability

- OAuth tokens encrypted at rest
 - Firebase JWT verification on backend
 - Least-privilege scopes
 - Idempotent posting
 - Retry with backoff
 - Dead-letter queues
 - Secrets via GCP Secret Manager
-

13. Cost Controls

- Serverless compute for non-GPU workloads
- Single GPU inference service
- Quantized model
- Request batching

- Autonomous mode volume caps
 - Rate limiting per user
-

14. Deliverables Summary

This document defines:

- What will be built
- How it will be built
- In what order it will be delivered
- How safety, cost, and trust are enforced

Any changes require a formal amendment.

15. Approval

Proceeding with implementation confirms alignment with this technical specification.